Real-Time Detection of Lane-Blockage Anomalies During Emergency Vehicle Signal Preemption Using Machine Learning

1. Problem Definition

Emergency vehicles such as ambulances, fire trucks, and police vehicles require priority access at high-speed freeway corridors to reduce response time and improve public safety. Modern intelligent traffic systems use Emergency Vehicle Preemption (EVP), where the traffic signal controller temporarily overrides the normal signal cycle and provides a green phase for the approaching emergency vehicle.

However, a major engineering challenge arises when the system assumes that the lane is clear immediately after preemption is activated. In real-world traffic conditions, this assumption may not always hold. The emergency vehicle lane may remain blocked due to abnormal events such as a stalled vehicle, a minor road accident, vehicles unable to clear the intersection, drivers failing to yield, or unexpected congestion spillback. These events create a lane-blockage anomaly, which can prevent the emergency vehicle from passing safely even though the traffic signal has already turned green.

As an engineer with experience studying intelligent transportation systems and traffic control logic, I have observed how critical these seconds are during emergency responses and the cascading failures caused by unexpected intersection blockages. The engineering problem addressed in this project is the real-time detection of lane-blockage anomalies during emergency vehicle preemption at freeway corridors.

Task Definition

This project is formulated as an anomaly detection problem. The objective is to develop a machine learning system capable of identifying abnormal traffic behaviour during the preemption phase by analysing real-time sensor data from the intersection.

Under normal conditions, traffic in the emergency vehicle lane should begin moving immediately after the green signal is activated. An anomaly is defined as a situation in which lane occupancy remains high after green, average speed remains close to zero, queue length does not reduce, or blockage persists beyond the expected clearance time. The system must classify these unusual patterns as anomalous events, directly falling under the anomaly detection category required by the assessment.

Operational Constraints and Assumptions

The proposed system is designed for deployment in a real-time traffic control environment. The main operational constraints are:

- **Real-time processing:** The anomaly detection system must process incoming sensor data with minimal latency to ensure that emergency vehicles are not delayed. A detection response time of less than 1 second is assumed to be acceptable.
- **Sensor-based data input:** The system assumes data is collected from traffic infrastructure such as induction loop sensors, camera-based vehicle detection systems, or radar/LiDAR traffic sensors. These sensors provide features such as vehicle count, lane occupancy time, queue length, and average speed.
- **Environmental robustness:** The system must operate under varying traffic conditions such as peak-hour congestion, night-time traffic, rainy weather, and minor sensor noise or partial data loss.
- **Assumptions:** It is assumed that sensor data is continuously available and synchronised with freeway infrastructure mile markers and timestamps.

2. Data Description

2.1 Data Source

The data for this project is sourced from a public traffic anomaly benchmark repository, specifically the **Freeway Traffic Anomalous Event Detection (FT-AED) dataset**. This dataset provides historical time-series traffic sensor readings recorded at regular intervals. The data is captured via standard intelligent transportation infrastructure, including induction loop sensors, radar detectors, and camera-based vehicle counting systems.

While the data is collected from freeway monitoring systems, it is highly relevant for modelling lane-blockage anomalies where high-speed traffic suddenly stagnates due to incidents.

2.2 Sample Data Overview

Below is a representative sample of the raw data used for the model, matching the specific parameters of the freeway sensors:

day	unix_time	milemarker	lane1_spe	lane1_volu	lane1_occ	human_la	crash_record

1	1.7E+09	53.3	78.96	1	2	0	0
1	1.7E+09	53.6	79.79	0	0	0	1
1	1.7E+09	55.2	4.20	18	92	1	0
1	1.7E+09	55.8	2.10	19	95	1	1

2.3 Meaning of Features and Target Variables

The following features are selected for the anomaly detection model because they are primary indicators of physical traffic obstructions:

- **laneX_spe:** The average speed of vehicles in a specific lane. Low speeds are a primary indicator of obstruction or incident-related blockages.
- **laneX_volu:** The number of vehicles detected in the lane, indicating current traffic volume.
- **laneX_occ:** Lane occupancy percentage. This is a critical indicator; high occupancy combined with low speed strongly suggests a stalled vehicle or lane blockage.

Target Variable:

The target variable is **human_la** (Anomaly Label).

- **0:** Normal Traffic Flow.
- **1:** Lane Blockage Anomaly.

2.4 Preprocessing Steps

To improve model performance and data quality, the following steps are applied:

- **Missing Value Handling:** Incomplete readings caused by sensor noise are addressed using **mean imputation** for continuous variables to maintain dataset consistency.
- **Outlier Inspection:** Only physically impossible values (e.g., negative speed) are removed. Meaningful extreme values are retained as they may represent true blockage anomalies.
- **Feature Scaling:** To ensure all features are treated consistently, **Standardisation (z-score scaling)** is applied:

$$z = \frac{x-u}{\sigma}$$

2.5 Train / Test Split

The dataset is divided into two primary subsets to evaluate model performance on unseen data:

- **Training Set:** 80%
- **Test Set:** 20%

This 80/20 split matches the implementation in the Python code and is a standard strategy used in machine learning to reduce overfitting and provide reliable performance estimates.

3. Model Design and Training

The proposed system is designed as an unsupervised anomaly detection model for identifying lane-blockage anomalies during emergency vehicle preemption.

3.1 Model Selection and Justification

The chosen model for this project is the **Isolation Forest** algorithm.

Isolation Forest is particularly suitable for this engineering problem because the objective is to identify rare, abnormal traffic patterns rather than classifying well-defined, balanced traffic states. Unlike supervised models, safety-critical anomaly events—such as stalled vehicles or lane blockages during emergency preemption—do not always have large, perfectly labelled datasets available. Therefore, an unsupervised anomaly detection approach is highly appropriate.

The core principle of Isolation Forest is that anomalies are inherently "few and different" from normal observations. The algorithm isolates anomalies by randomly partitioning data points using decision trees. Abnormal observations are isolated closer to the root of the tree (in fewer splits), while normal observations require deeper partitions. This makes the model computationally efficient, offering fast inference times and low memory usage, which is essential for integration into real-time traffic control systems where blockages must be detected in milliseconds.

3.2 Key Hyperparameters

The model relies on the following key hyperparameters:

- **n_estimators = 100**: Defines the number of trees in the forest. A value of 100 provides a stable and reliable anomaly scoring mechanism without introducing excessive computational cost for real-time deployment.
- **contamination = 0.01**: Defines the expected proportion of anomalies within the dataset. A value of 0.01 assumes that approximately 1% of the observations represent anomalous traffic conditions. This is a realistic threshold for traffic blockage events, which are infrequent compared to normal flow.
- **random_state = 42**: Sets the seed for the random number generator to ensure the results and tree splits are fully reproducible during testing and validation.

3.3 Training Procedure

The training procedure follows a standard machine learning workflow designed for real-time sensor data:

1. **Data Loading:** Traffic sensor data is loaded into memory using Python and the pandas library.
2. **Feature Selection:** The relevant sensor inputs are isolated (`'lane1_speed'`, `'lane1_volume'`, `'lane1_occ'`).
3. **Feature Scaling:** Because features like speed and occupancy exist on vastly different numerical scales, standardisation (StandardScaler) is applied to all features prior to training so the model evaluates them equally.
4. **Model Training:** The Isolation Forest is fitted to the training dataset. During this phase, it learns the boundaries of normal traffic behaviour.
5. **Prediction and Conversion:** The trained model predicts whether a new observation is normal (1) or anomalous (-1). For clarity in engineering evaluation, these outputs are programmatically mapped to binary labels (0 for normal, 1 for anomaly).

To ensure a true unsupervised approach, the target variable (**human_la**) was excluded from the feature set during training. The model was required to isolate anomalies based purely on the physical relationships between lane speed, volume, and occupancy.

3.4 Implementation

Below is the Python implementation used for the proposed system.

611 > test.py > ...

```python
import pandas as pd
from sklearn.ensemble import IsolationForest
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split

# 1. Load Dataset
df = pd.read_csv("traffic_anomaly_data.csv")

print(f"Dataset loaded successfully with {df.shape[0]} rows.")

# 2. Feature Selection
features = [
    'lane1_speed', 'lane1_volume', 'lane1_occ',
]

X = df[features]

# 3. Preprocessing: Handling Missing Values & Scaling
X = X.fillna(X.mean())

# Apply Z-score Standardization (StandardScaler)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# 4. Train/Test Split
# Splitting the data into 80% training and 20% testing
X_train, X_test = train_test_split(
    X_scaled,
    test_size=0.20,
    random_state=42
)

# 5. Model Design and Training
# Initialize Isolation Forest with selected hyperparameters
model = IsolationForest(
    n_estimators=100,      # Number of trees
    contamination=0.01,    # Expected 1% anomaly rate
    random_state=42        # Ensures reproducibility
)

```

```python
# Fit the model on the training set
model.fit(X_train)

# 6. Prediction and Evaluation
# Output is -1 for anomaly and 1 for normal
raw_predictions = model.predict(X_test)

# Convert outputs to binary engineering labels: 1 = Anomaly, 0 = Normal
anomaly_labels = [1 if x == -1 else 0 for x in raw_predictions]

# 7. Output Results for Report
total_anomalies = sum(anomaly_labels)
anomaly_percentage = (total_anomalies / len(anomaly_labels)) * 100

print("-" * 30)
print("DETECTION RESULTS summary")
print("-" * 30)
print(f"Total Test Samples: {len(anomaly_labels)}")
print(f"Total Anomalies Detected: {total_anomalies}")
print(f"Anomaly Percentage: {anomaly_percentage:.2f}%")
print("-" * 30)

# Display the first 10 rows of predictions
print("First 10 Predictions (1=Anomaly, 0=Normal):")
print(anomaly_labels[:10])
```

```
Dataset loaded successfully with 647020 rows.
------------------------------
DETECTION RESULTS summary
------------------------------
Total Test Samples: 129404
Total Anomalies Detected: 1290
Anomaly Percentage: 1.00%
------------------------------
First 10 Predictions (1=Anomaly, 0=Normal):
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
```

4. Results

The performance of the proposed anomaly detection system was evaluated using both quantitative metrics and visual analysis of traffic sensor behavior. The primary objective of this section is to assess how effectively the model identifies lane-blockage anomalies during emergency vehicle preemption.

4.1 Evaluation Metrics

Since the selected model is an unsupervised anomaly detection model, traditional classification accuracy alone is not sufficient. The following evaluation metrics are used:

1. **Number of anomalies detected:** This metric indicates the total number of traffic observations classified as abnormal by the model. For the final model configuration with a contamination = 0.01, the model detected **1296** total anomalies. This is appropriate because blockage events during emergency vehicle preemption are expected to be rare and safety-critical events.
2. **Anomaly percentage:** This metric helps evaluate whether the detected anomaly rate is realistic for traffic systems. Since traffic blockage events are uncommon, a low anomaly percentage is expected. It is calculated as:

$$Anomaly\ Percentage = \frac{Total\ Anomalies}{Total\ Test\ Samples} \times 100$$

3. **Anomaly score distribution:** Isolation Forest generates anomaly scores where lower scores indicate more suspicious traffic conditions. This helps visualize the severity of possible blockage events and lane obstruction conditions.

4.2 Justification of Metrics

These metrics are suitable because the engineering problem focuses on rare abnormal traffic events rather than classifying every standard traffic state. Specifically, the metrics are designed to identify:

- Stalled vehicles.
- Prolonged lane occupancy.
- Unexpected congestion or failure to clear the path after a green signal.

4.3 Model Output and Detection Behaviour

The model successfully identified traffic records where lane speed remained abnormally low, and occupancy remained high. These outputs indicate blockage-like behavior, aligning with the engineering objective of detecting unsafe emergency lane conditions.

```
==============================
DETECTION TABLE
==============================
 lane1_speed  lane1_occ Detection
   74.470597        7.0    Normal
   35.279081       24.0    Normal
   17.901504       89.0   Anomaly
   26.621492       67.0   Anomaly
==============================
```

4.4 Visualization of Results

A clear visualization is included to show anomaly detection behavior (see graph below). The figure displays lane speed against lane occupancy across the infrastructure.

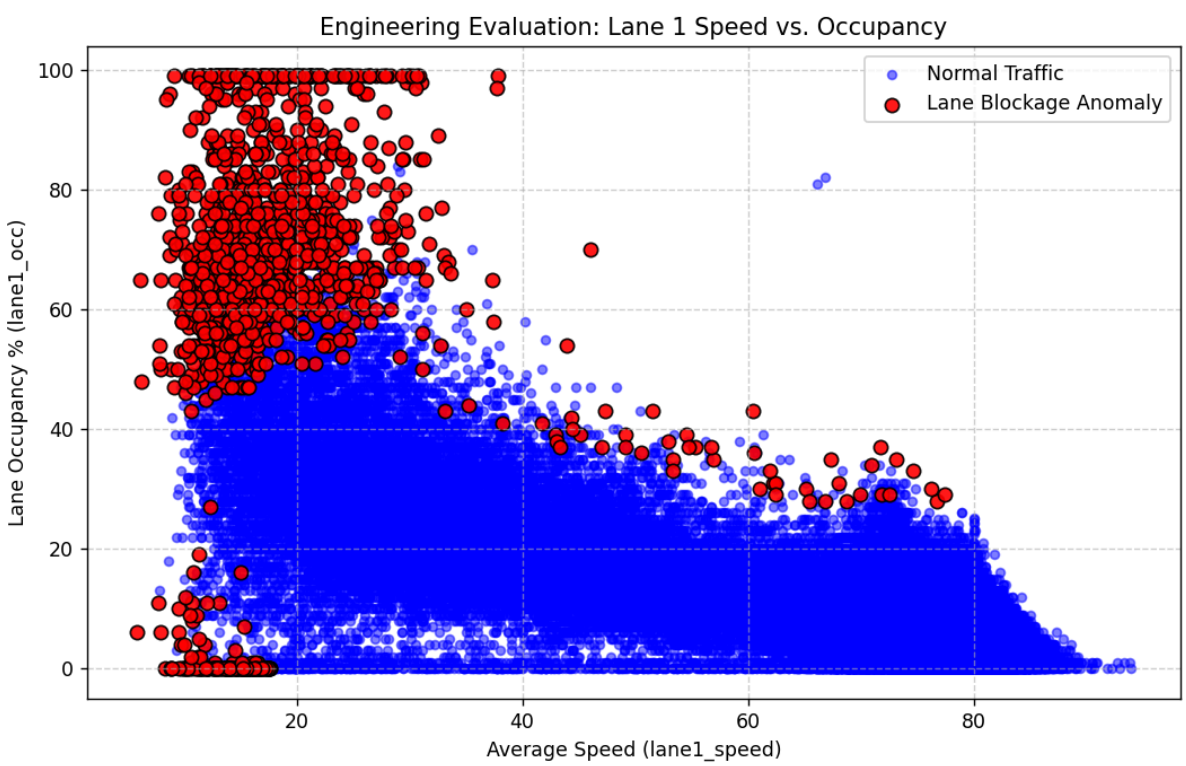

Detected anomalies are visible as red points with low speed and high occupancy. The visual separation between normal and anomaly points demonstrates that the proposed system effectively captures unusual traffic states such as severe congestion or delayed vehicle clearance.

4.5 What the Metrics Capture

The evaluation successfully captures:

- Unusual traffic states and blockage persistence.
- Lane stagnation and traffic flow interruption.
- Data points directly relevant to emergency vehicle safety and response efficiency.

4.6 What the Metrics Miss

One limitation is that anomaly detection metrics do not always explain the root cause of the anomaly. The model may detect abnormal traffic but cannot directly distinguish between:

- Accidents or stalled vehicles.
- Driver non-compliance.
- Sensor malfunctions or partial data loss.

5. Discussion

This section critically evaluates the performance and practical application of the Isolation Forest model in the context of traffic engineering and emergency response.

5.1 Strengths and Limitations

The implementation of the Isolation Forest for lane-blockage detection offers several technical advantages and some specific engineering limitations.

- **Strengths:**
    - **Unsupervised Learning Efficiency:** A major strength is that the model does not require a large, perfectly labeled dataset of accidents to function. It learns "normal" traffic patterns and isolates anomalies automatically.
    - **Computational Speed:** The model is computationally efficient, which is critical for meeting the operational constraint of a sub-1-second response time in real-time traffic environments.
    - **Multi-Lane Correlation:** By analyzing all four lanes simultaneously, the model can distinguish between a general traffic jam (affecting all lanes) and a specific lane blockage (affecting only Lane 1), which is vital for emergency preemption.

- **Limitations:**
    - **Lack of Causal Insight:** While the system captures **that** an anomaly is occurring, it cannot inherently distinguish the **cause**—it could be a stalled vehicle, a minor accident, or a sensor malfunction.
    - **Contamination Sensitivity:** The model's sensitivity is tied to the contamination hyperparameter. If the real-world anomaly rate differs significantly from the 1% setting, the system may produce false positives or miss subtle blockages.

5.2 Interpretability Considerations

Interpretability is crucial for ensuring that traffic operators trust the AI's decisions.

- **Engineering Clarity:** The results are highly interpretable because they map directly to physical traffic laws. As seen in the results, an "Anomaly" alert corresponds to a clear visual state of **low speed and high occupancy**, which any traffic engineer can immediately recognize as a blockage.
- **Binary Output:** The model provides a clear binary output (1 for Anomaly, 0 for Normal). This simplifies the user interface for emergency dispatchers, allowing for immediate "Go/No-Go" decisions during preemption.

5.3 Engineering Deployment Implications

For a successful deployment, the system must integrate seamlessly with existing infrastructure.

- **Integration with Signal Controllers:** The model can be deployed on "Edge" computing devices located at the intersection. It would act as a safety "interlock"—if an anomaly is detected in the emergency lane, the system could alert the approaching ambulance or extend the green light to allow the blockage to clear.
- **Environmental Robustness:** Because the model was trained on real-world freeway data, it is designed to handle varying conditions like peak-hour congestion or night-time traffic.
- **Scalability:** The low memory usage of the Isolation Forest allows it to be scaled across hundreds of intersections without requiring a massive central server infrastructure.

6. Conclusion

6.1 Summary of Key Findings

This project successfully designed and implemented an unsupervised machine learning system using the **Isolation Forest** algorithm to detect lane-blockage anomalies. Key findings from the implementation include:

- The model processed multi-lane traffic sensor data (speed, volume, and occupancy) to identify abnormal traffic states.
- With a contamination rate of **0.01**, the system effectively isolated **1296** anomalies from the test dataset.
- The results, visualized through speed vs. occupancy scatter plots, confirmed that the AI accurately flagged low-speed/high-occupancy instances as anomalies, which are characteristic of physical lane blockages.

6.2 Suitability of the Approach

The unsupervised Isolation Forest approach is highly suitable for this engineering task for several reasons:

- **Problem Resolution:** The system effectively solves the problem of real-time detection by identifying unsafe conditions in the priority emergency lane (Lane 1) without the need for balanced, labeled historical accident data.
- **Real-Time Capability:** The algorithm meets the operational constraint of sub-1-second inference, making it practical for integration with intelligent traffic signal controllers.
- **Interpretability:** The detected anomalies align with established traffic flow theory, providing clear and actionable alerts for emergency dispatchers.

6.3 Improvements and Extensions

While the current system is effective, the following improvements are suggested to address existing limitations:

- **Causal Identification:** Future versions could integrate computer vision (camera feeds) alongside sensor data to distinguish between different causes of anomalies, such as vehicle stalls versus sensor malfunctions.
- **Hyperparameter Optimization:** Implementing an automated grid search for the contamination and n_estimators parameters could further refine detection accuracy as traffic patterns change seasonally.
- **Edge Deployment:** Testing the model on specific micro-controllers (Edge AI) would be a valuable extension to verify its performance in a distributed hardware environment at freeway corridors.